
MGR-SID: HIERARCHY-AWARE GRAPH-REGULARIZED SEMANTIC IDS FOR GENERATIVE RECOMMENDATION

Anonymous authors

Paper under double-blind review

ABSTRACT

Generative recommenders depend heavily on the quality of the item tokenizer that maps each item to a discrete semantic ID. Recent collaborative tokenizer designs usually inject one globally fused collaborative representation into the entire tokenizer, but our study of MiniOneRec suggests that this framing is too coarse for hierarchical semantic IDs. We find that the dominant tokenizer bottleneck is not full-code collision after generation, but local ambiguity inside semantically plausible prefixes, where the prefix is correct while the final leaf token remains hard to predict. Based on this observation, we propose **MGR-SID**, a hierarchy-aware graph-regularized tokenizer built on top of MiniOneRec’s semantic RQ-VAE backbone. The current **v1** instantiation uses a train-only graph bank with a purified coarse collaborative graph, a spectral mid-resolution graph, and a purified local transition graph, and applies different graph regularizers to different SID levels. We first show on Industrial and Office diagnostic probes that spectral middle-resolution graph views are much stronger than heuristic middle views for same-prefix disambiguation. We then align the tokenizer training regime with the upstream MiniOneRec recipe and evaluate training-time graph regularization under this fair setting. On Industrial, the hierarchy-aware variant reduces the final generated index collision from 0.00353 to 0.00326 relative to the semantic baseline, and decreases the fraction of test targets in multi-leaf same- l_2 buckets from 68.94% to 61.31%. These results suggest that graph-structured collaboration is most useful when it enters SID learning as level-specific structural supervision rather than uniform feature fusion.

1 INTRODUCTION

Generative recommendation replaces ranking over a catalog with sequence generation over item codes. This reformulation has become increasingly attractive because it allows recommender systems to inherit the training and decoding machinery of language models. Recent systems such as TIGER, LCRec, and MiniOneRec demonstrate that a discrete item tokenizer can be the bridge between item semantics and autoregressive recommendation (Rajput et al., 2023; Zheng et al., 2023; Kong et al., 2025). Once the tokenizer is fixed, however, its structural biases are inherited by every downstream stage, including constrained decoding, supervised fine-tuning, and recommendation-oriented reinforcement learning.

This makes tokenizer quality a first-order problem rather than a preprocessing detail. In our current MiniOneRec reproduction line, we repeatedly observed a pattern that does not fit the usual “collision is the core problem” narrative. After the official `sid-generate` collision-repair stage, the final code collision rate is already at the sub-percent level, yet many prediction errors remain. A closer look shows that these errors often occur after the model has already entered the correct semantic region. In other words, the prefix is frequently correct, but the final leaf token is not. This points to a more specific bottleneck: *local ambiguity inside semantically plausible subtrees*.

At the same time, existing collaborative tokenizer methods suggest that collaborative signals should influence item tokenization, but they often do so through a single globally fused collaborative representation (Fang et al., 2026; Liang et al., 2026; Wang et al., 2026; Liu et al., 2025). That design

is natural for flat item representations, but it is less natural for hierarchical semantic IDs. Different levels in a semantic ID serve different structural roles. Upper levels act as coarse routing decisions, while deeper levels distinguish among increasingly similar candidates. A single collaborative view injected uniformly across all levels may therefore be both under-specified and unstable.

This paper develops the current state of that idea into a concrete tokenizer-side method. We present **MGR-SID**, a hierarchy-aware graph-regularized semantic ID framework that uses graph structure as the carrier of collaborative information. Rather than replacing the semantic tokenizer, MGR-SID keeps MiniOneRec’s semantic RQ-VAE backbone and augments it with three train-only graph views: a purified coarse collaborative graph, a middle-resolution spectral graph, and a purified local transition graph. In the current **v1** implementation, we instantiate hierarchy-awareness with a simple but explicit level-wise rule: coarse regularization at level 1, spectral middle-view regularization at level 2, and local transition regularization at level 3.

Our evidence proceeds in three steps. First, we run diagnostic probes on Industrial and Office and show that middle-resolution graph views are indispensable: the best spectral middle view substantially outperforms heuristic middle views on same- l_2 ambiguity buckets. Second, we trace the provenance of the original MiniOneRec tokenizer and show that fair tokenizer comparison requires the upstream RQ-VAE training regime rather than the default local configuration. Third, under this aligned training setup, we show that hierarchy-aware graph regularization improves the *final generated SID* even though it does not dominate raw train-stage collision.

The current draft is intentionally scoped to the tokenizer stage. We do not yet claim downstream SFT or RL gains. Instead, the goal is to establish a precise tokenizer-side story before carrying the new SID into the rest of the pipeline.

Contributions. This draft makes three contributions.

1. We identify a tokenizer-side bottleneck for MiniOneRec: after fair reproduction of the upstream pipeline, the remaining difficulty lies in local leaf ambiguity rather than in unrepaired global collision.
2. We propose MGR-SID, a hierarchy-aware graph-regularized semantic ID framework that injects graph-structured collaboration into a semantic tokenizer through level-specific structural constraints instead of uniform feature fusion.
3. We provide preliminary but concrete evidence for this design. On Industrial, under upstream-aligned RQ-VAE training, hierarchy-aware regularization reduces final SID collision from 0.00353 to 0.00326 and improves local same- l_2 ambiguity metrics, while uniform graph regularization is clearly harmful.

The remainder of the paper is organized as follows. Section 2 situates MGR-SID in prior work. Section 3 summarizes the tokenizer pipeline and the empirical bottlenecks that motivate our design. Section 4 presents the current MGR-SID v1 method. Section 5 reports probe studies, provenance alignment, and tokenizer-stage results. Section 6 concludes with limitations and the next steps toward downstream SFT evaluation.

2 RELATED WORK

Generative recommendation and semantic item IDs. Recent generative recommenders reformulate recommendation as constrained sequence generation over discrete item tokens. TIGER introduced generative retrieval for recommendation and showed that autoregressive decoding over item codes can replace traditional ranking (Rajput et al., 2023). LCRec aligned collaborative semantics with language-model adaptation and demonstrated that tokenized item representations can be integrated into large language model pipelines (Zheng et al., 2023). MiniOneRec pushed this line further by releasing an end-to-end open-source framework that covers semantic ID construction, SFT, RL, and evaluation (Kong et al., 2025). These works establish the importance of tokenizer quality, but they mostly inherit a semantic-first tokenizer view. Our work stays within the MiniOneRec pipeline and asks a narrower question: how should collaborative structure enter a *hierarchical* semantic tokenizer without destabilizing it?

Collaborative-aware tokenization for generative recommenders. Several recent papers argue that text-only tokenization is insufficient. ReSID motivates recommendation-native tokenizer design from the perspective of prefix-conditional predictability and semantic quantization beyond plain language modeling (Liang et al., 2026). PRISM introduces purified collaborative representations and semantic anchoring to avoid noisy collaborative tokenization (Fang et al., 2026). PIT further highlights the volatility of collaborative signals and studies dynamic personalized tokenization (Wang et al., 2026). DiscRec separates semantic and collaborative factors before fusion (Liu et al., 2025). These methods all strengthen the case that collaboration matters, but they still leave open the question that matters most for hierarchical IDs: should the same collaborative signal be reused across every SID level, or should different levels preserve different structures? MGR-SID is built precisely around that missing allocation question.

Tokenization objectives beyond reconstruction. Another relevant thread asks tokenizers to preserve recommendation structure rather than optimizing reconstruction alone. ETEGRec argues that tokenization should be aligned with the recommender objective itself instead of being treated as an isolated preprocessor (Liu et al., 2024). CoST uses contrastive quantization to preserve semantic tokenization quality for generative recommendation (Zhu et al., 2024). These works are close in spirit to our method because they move the burden of recommendation structure into tokenizer learning. Our difference is that we use graph-structured collaborative signals as explicit structural targets and distribute them across SID levels, instead of introducing a single tokenization objective that treats all relations uniformly.

Graph signal processing and transition-aware recommendation. The graph view bank in MGR-SID is inspired by work showing that different graph scales capture different recommendation patterns. GSPRec decomposes graph signals into different spectral components for recommendation (Rabiah & McAuley, 2025), while FaGSP studies frequency-aware graph signal processing for collaborative filtering (Xia et al., 2024). Collaboration and Transition emphasizes that long-range collaborative structure and short-range transition structure are complementary rather than interchangeable (Zhu et al., 2023). These papers are not SID tokenizers, but they provide a strong clue for our design: the collaborative information most useful for a hierarchical tokenizer is likely multi-resolution, and middle-scale graph structure deserves explicit modeling rather than being approximated by a crude interpolation between global and local views.

Positioning of this work. MGR-SID combines ideas from all three threads above, but it does not aim to compete as a generic graph encoder or a full end-to-end personalized tokenizer. The current v1 is deliberately minimal. It inherits MiniOneRec’s semantic RQ-VAE backbone, borrows the denoising discipline of PRISM, uses the predictability-based motivation of ReSID, and instantiates graph-aware supervision with a coarse/mid/local graph bank whose middle view is selected by empirical probes inspired by GSPRec and FaGSP. In this sense, the paper’s main point is not “better graph features,” but *level-specific graph supervision for semantic ID learning*.

3 BACKGROUND AND PROBLEM FORMULATION

3.1 MINIONEREC TOKENIZER PIPELINE

Let each item i be associated with a semantic embedding $\mathbf{x}_i \in \mathbb{R}^d$. MiniOneRec tokenizes items by training an RQ-VAE over these item embeddings and then decoding each item into a fixed-length semantic ID $\text{SID}(i) = [z_i^{(1)}, z_i^{(2)}, z_i^{(3)}]$ (Kong et al., 2025; Lee et al., 2022). During `sid-train`, the model learns residual quantizers over the semantic embedding space. During `sid-generate`, the model re-encodes colliding items with Sinkhorn-based assignment enabled so that the final exported `index.json` has much lower collision than the raw train-stage codes.

This separation is important. A tokenizer should not be judged only by train-stage collision. The actual pipeline uses `sid-generate` as part of the final tokenizer construction, so a fair comparison must examine the final generated SID rather than raw codes emitted directly by the training loop.

Table 1: Tokenizer provenance alignment on Industrial. The upstream-style RQ-VAE regime reproduces the low-collision official tokenizer quality, whereas the local default training configuration does not.

Setting	Epochs	Batch size	Best train collision	Final index collision
Local default rerun	10	256	0.9989	0.9940
Upstream-style rerun	10000	20480	0.0993	0.00434
Official committed index	–	–	–	0.00434

3.2 WHY LOCAL AMBIGUITY MATTERS MORE THAN UNREPAIRED GLOBAL COLLISION

Our diagnosis of the current semantic baseline consistently points to a local failure mode. We observed many cases where the model reaches the correct semantic prefix but fails at the final leaf. This means that the tokenizer still groups together items that are semantically close yet behaviorally distinct. We refer to this phenomenon as *local leaf ambiguity*.

To make this concrete, consider an item’s first two SID levels $(z_i^{(1)}, z_i^{(2)})$. The set of items sharing this two-level prefix defines a same- l_2 bucket. If many leaves remain under the same prefix, then the final decoder still faces a difficult local choice even when the upper routing is already correct. Our current analysis therefore treats same- l_2 crowding and the conditional entropy of the third-level code given the prefix as tokenizer health indicators in addition to collision.

3.3 FAIR TOKENIZER COMPARISON REQUIRES UPSTREAM-ALIGNED TRAINING

Before testing graph-aware tokenizers, we first had to resolve a provenance problem in the reproduced MiniOneRec pipeline. A direct rerun with the current local default configuration produced extremely high train-stage and generate-stage collision, which appeared inconsistent with the official low-collision Industrial index already stored in the repository. By comparing the local code path with the upstream MiniOneRec release and then rerunning with the upstream tokenizer recipe, we found that the discrepancy came primarily from hyperparameters rather than a broken implementation.

Table 1 is central to the rest of the paper. It shows that fair tokenizer comparison must use the upstream-style RQ-VAE training regime. All training-time comparisons for MGR-SID in this draft therefore adopt the same tokenizer backbone, optimizer, and training schedule as the upstream MiniOneRec recipe.

3.4 PROBLEM STATEMENT

These observations motivate a more specific tokenizer question than “should collaboration be fused into tokenization?” We instead ask:

How should graph-structured collaborative information be injected into different SID levels so that a semantic tokenizer becomes less locally ambiguous without collapsing its hierarchy?

The current MGR-SID v1 instantiation answers this question conservatively. It preserves the semantic RQ-VAE backbone and adds graph-based regularization only as a structural training signal. This allows us to test the value of hierarchy-aware graph supervision before modifying the downstream SFT or RL stages.

4 METHOD

4.1 OVERVIEW

MGR-SID keeps MiniOneRec’s semantic RQ-VAE tokenizer and augments it with a train-only graph bank. The core intuition is simple: different SID levels solve different structural discrimination problems, so they should preserve different collaborative graph structures. In the current **v1** implementation, we instantiate this idea with a fixed but non-uniform level-to-view assignment:

- level 1 is regularized by a coarse collaborative graph;
- level 2 is regularized by a middle-resolution spectral graph;
- level 3 is regularized by a local transition graph.

This is intentionally simpler than the full long-term vision of learned level-wise allocation. The purpose of v1 is to test whether hierarchy-aware structural supervision is already useful before adding another learned allocation module.

4.2 SEMANTIC BACKBONE

Given an item embedding $\mathbf{x}_i$, the encoder produces a latent representation that is quantized by three residual quantizers. Let $\mathbf{q}_i^{(l)}$ be the quantized vector produced by the l -th quantizer, and let

$$\mathbf{h}_i^{(l)} = \sum_{r=1}^l \mathbf{q}_i^{(r)} \quad (1)$$

be the cumulative quantized representation after level l . The decoder reconstructs $\mathbf{x}_i$ from $\mathbf{h}_i^{(3)}$. This is exactly the same semantic tokenizer backbone used by MiniOneRec, and we keep its reconstruction loss and residual quantization loss unchanged.

4.3 GRAPH BANK CONSTRUCTION

All graph views are constructed from training interactions only. For each interaction sample with history $H_t = [i_{t-k}, \dots, i_{t-1}]$ and target item y_t , we build three graph carriers.

Coarse collaborative graph. We build an undirected item-item graph from the unique items appearing in the recent history window and the target. For items u and v that co-occur inside the sequence, we add a symmetric edge weighted by inverse distance. This graph captures broad compatibility and stable collaborative organization. We then purify it by support pruning, popularity debiasing, and row normalization:

$$\mathbf{A}_{\text{coarse}} = \text{RowNorm}(\text{Debias}(\text{Prune}(\mathcal{G}_{\text{coarse}}))) . \quad (2)$$

In the current implementation, pruning removes edges with weight below 2.0 and popularity debiasing uses exponent $\alpha = 0.5$.

Local transition graph. We build a directed graph from recent history items to the target item. Each edge receives inverse-rank weight, so more recent items contribute more. The purified local graph applies support pruning, target-popularity correction, and row normalization:

$$\mathbf{A}_{\text{local}} = \text{RowNorm}(\text{TargetDebias}(\text{Prune}(\mathcal{G}_{\text{local}}))) . \quad (3)$$

This view captures short-range next-item preference but is naturally sparse and noisy.

Middle-resolution graph. The middle view is the key design choice. Early heuristic probes using diffusion residuals and simple band-pass proxies already suggested that mid-scale structure matters, but they were weaker than spectral graph constructions. The current v1 therefore selects a FaGSP-inspired spectral middle view as the default:

$$\mathbf{A}_{\text{mid}} = \text{BandReconstruct}(\text{SymNorm}(\mathbf{A}_{\text{coarse}})) . \quad (4)$$

Operationally, we compute a band-limited spectral reconstruction of the purified coarse graph by retaining eigen-components in a middle spectral range. The implementation uses a rank of 48 and an eigen-ratio range of $[0.25, 0.65]$. We also explored a GSPRec-style temporal spectral view as an alternative, but use the FaGSP-style view as the main G_{mid} because it performed best in our diagnostic probes.

Semantic-anchor purification. Inspired by PRISM, we additionally tested a semantic-anchor purification stage in which a purified graph is intersected with a top- k semantic k NN graph extracted from the text embedding space. This module improves some graph views, but our probes indicate that the main gain comes from the quality of the middle-resolution graph itself rather than from anchor-only denoising. In the current training-time v1, the final chosen middle view is the spectral graph built from the purified coarse graph.

4.4 HIERARCHY-AWARE GRAPH REGULARIZATION

The semantic backbone already optimizes

$$\mathcal{L}_{\text{rec}} + \mathcal{L}_{\text{rq}}, \quad (5)$$

where $\mathcal{L}_{\text{rec}}$ is the reconstruction loss and $\mathcal{L}_{\text{rq}}$ is the residual quantization loss. MGR-SID v1 adds graph regularization on the cumulative representations $\mathbf{h}_i^{(l)}$.

For a mini-batch $\mathcal{B}$ and a row-normalized subgraph adjacency $\mathbf{A}_{\mathcal{B}}^{(l)}$ extracted from the selected graph view, we define

$$\mathcal{L}_{\text{graph}}^{(l)} = \frac{1}{|\mathcal{B}|} \sum_{i \in \mathcal{B}} \left\| \mathbf{h}_i^{(l)} - \sum_{j \in \mathcal{B}} \mathbf{A}_{\mathcal{B}, ij}^{(l)} \mathbf{h}_j^{(l)} \right\|_2^2. \quad (6)$$

This loss favors level- l representations that are locally smooth on the corresponding graph view. Intuitively, items that are connected in the relevant graph should remain nearby at the level that is responsible for preserving that structure.

The hierarchy-aware v1 objective is then

$$\mathcal{L}_{\text{total}} = \mathcal{L}_{\text{rec}} + \mathcal{L}_{\text{rq}} + \lambda_1 \mathcal{L}_{\text{graph}}^{(1)} + \lambda_2 \mathcal{L}_{\text{graph}}^{(2)} + \lambda_3 \mathcal{L}_{\text{graph}}^{(3)}, \quad (7)$$

where

$$\mathbf{A}_{\mathcal{B}}^{(1)} \leftarrow \mathbf{A}_{\text{coarse}}, \quad \mathbf{A}_{\mathcal{B}}^{(2)} \leftarrow \mathbf{A}_{\text{mid}}, \quad \mathbf{A}_{\mathcal{B}}^{(3)} \leftarrow \mathbf{A}_{\text{local}}. \quad (8)$$

In the current implementation, $(\lambda_1, \lambda_2, \lambda_3) = (0.05, 0.15, 0.05)$.

4.5 CONTROLS

To isolate the effect of hierarchy-awareness, we compare MGR-SID v1 with two controls:

- **semantic baseline:** the upstream-aligned MiniOneRec RQ-VAE tokenizer with no graph regularization;
- **uniform regularization:** a control that averages the coarse, middle, and local graphs and applies one graph smoothness loss to the final cumulative representation only.

The uniform control keeps the same graph ingredients but removes the hierarchy-aware assignment. It is therefore the most direct test of whether the benefit comes from non-uniform level-specific supervision rather than from adding graph information in any form.

4.6 CURRENT SCOPE

MGR-SID v1 should be read as a tokenizer-stage instantiation, not the final method ceiling. It does *not* yet learn the graph allocation itself, and it does *not* modify downstream SFT or RL. The purpose of the current draft is narrower: to establish whether graph-structured collaborative information helps the semantic tokenizer *when injected as level-specific structural supervision*.

5 EXPERIMENTS

5.1 EXPERIMENTAL SETUP

Datasets. We use the Industrial and Office Amazon benchmarks already prepared in the MiniOneRec reproduction pipeline. Industrial contains 36,259 training interactions, 4,533 test interactions, and 3,686 items. Office contains 38,924 training interactions, 4,866 test interactions, and 3,459 items. Both use the repository’s existing time-based splits.

Tokenizer backbone. Unless otherwise noted, all training-time tokenizer comparisons use the upstream-aligned MiniOneRec RQ-VAE regime established in Table 1: three quantizers with codebook size 256, embedding dimension 32, AdamW, learning rate 10^{-3} , 10,000 epochs, and batch size 20,480.

Table 2: Best middle-view probe results on same- l_2 buckets. Spectral middle views are much stronger than the first heuristic middle views, which motivated the use of a spectral G_{mid} in MGR-SID v1.

Dataset	Best heuristic middle view	Best spectral middle view	Spectral view name
Industrial	0.2747	0.5401	fagsp_mid_base
Office	0.2381	0.7619	fagsp_mid_base

Table 3: Industrial tokenizer results under the upstream-aligned MiniOneRec training regime. Train-stage collision alone is misleading: the hierarchy-aware variant becomes the best tokenizer after sid-generate.

Method	Best train collision	Final index collision	Max conflict
Semantic baseline	0.1066	0.00353	2
Uniform graph regularization	0.2458	0.00597	6
Hierarchy-aware MGR-SID v1	0.1319	0.00326	2

Evaluation metrics. We report three classes of metrics.

- **Probe metrics:** target-better rate on ambiguity buckets, especially same- l_2 .
- **Tokenizer metrics:** best train-stage collision, final sid-generate collision, and maximum conflict size.
- **Local ambiguity metrics:** mean target l_2 leaf count, the fraction of targets in multi-leaf same- l_2 buckets, the fraction of targets in deep crowded same- l_2 buckets, and conditional entropy of level 3 under a fixed level-1/2 prefix.

5.2 DESIGN-SPACE PROBES: SPECTRAL MIDDLE VIEWS MATTER

We began with a non-invasive probe protocol that does not modify the tokenizer. Given an existing top-1 SID error, we ask whether a graph view scores the ground-truth target item above the currently favored wrong candidate. This lets us study which graph view is most informative before injecting it into training.

Table 2 provides the first direct justification for the graph bank design. The diagnostic result is not just that graph signals help. Rather, the result is that *middle-resolution graph structure matters enough to deserve explicit design*. The same- l_2 gain from the best spectral middle view is nearly doubled on Industrial and more than tripled on Office relative to the first heuristic middle view. This is what led us to instantiate G_{mid} with a spectral graph rather than a crude diffusion proxy.

5.3 TRAINING-TIME COMPARISON UNDER UPSTREAM-ALIGNED TOKENIZER TRAINING

We next test whether the graph signal survives once it is injected into tokenizer learning. All variants here use the fair upstream-aligned training setup from Table 1. We compare the semantic baseline, the uniform regularization control, and the hierarchy-aware MGR-SID v1 regularizer.

Table 3 reveals two important facts. First, uniform graph regularization is clearly harmful. This matters because it shows that adding graph structure per se is not enough. Second, train-stage collision and final tokenizer quality are not aligned. The semantic baseline still has the best raw train-stage collision, but after the official sid-generate stage, the hierarchy-aware tokenizer becomes the best final index. This is exactly why our evaluation protocol treats the final exported SID as the primary tokenizer artifact.

5.4 DOES THE FINAL TOKENIZER ACTUALLY REDUCE LOCAL AMBIGUITY?

The main question is whether the better final collision reflects the failure mode we care about. To answer this, we compare the final baseline index and the final hierarchy-aware index at the level of same- l_2 crowding.

Table 4: Local ambiguity comparison on Industrial final indices. The hierarchy-aware tokenizer reduces both catalog-level and test-weighted same- l_2 crowding.

Metric	Baseline	Hierarchy	Delta
Mean target l_2 leaf count	4.7999	4.4498	-0.3501
Fraction targets in multi-leaf l_2	0.6894	0.6131	-0.0763
Fraction targets in l_2 with ≥ 4 leaves	0.3283	0.2828	-0.0454
Mean target level-3 entropy under l_2	1.4533	1.2935	-0.1598

Table 4 shows that the improvement is structurally meaningful. The hierarchy-aware tokenizer does not merely shuffle code assignments while preserving the same ambiguity. It lowers the average number of competing leaves under the target prefix, reduces the fraction of targets that fall into multi-leaf same- l_2 buckets, and lowers the conditional entropy of the third-level code. In the full movement analysis, 17.32% of test targets move from a multi-leaf same- l_2 bucket into a singleton bucket, while only 9.68% move in the opposite direction. This strongly supports our original diagnosis that the remaining tokenizer bottleneck is local leaf ambiguity.

5.5 CURRENT INTERPRETATION

Taken together, the experiments support a cautious but coherent claim. We do *not* yet show downstream recommendation gains after SFT or RL. We do show that:

1. fair tokenizer comparison requires the upstream MiniOneRec training regime;
2. graph-structured collaborative information has real signal for local same- l_2 disambiguation;
3. this signal is strongest when encoded by a spectral middle-resolution graph;
4. hierarchy-aware graph regularization improves the *final generated SID* and makes the tokenizer less locally ambiguous than a fair semantic baseline.

For the current stage of the project, this is exactly the evidence we need before pushing the new SID into SFT. It validates the tokenizer-side hypothesis while keeping the scope honest.

6 CONCLUSION AND LIMITATIONS

This draft studies the tokenizer stage of MiniOneRec and argues that graph-structured collaboration should enter semantic ID learning as level-specific structural supervision rather than uniform feature fusion. The resulting method, MGR-SID v1, keeps the semantic RQ-VAE backbone intact, builds a train-only graph bank with coarse, middle, and local views, and applies hierarchy-aware graph regularization to different SID levels. Under a fair upstream-aligned tokenizer regime, this design improves the final generated Industrial SID over the semantic baseline and reduces same- l_2 local ambiguity.

The current evidence is still limited in three ways. First, the full training-time result is currently based on Industrial only. Second, the current v1 uses a fixed level-to-view assignment rather than a learned allocation module. Third, we have not yet carried the improved tokenizer into SFT, RL, and end-to-end recommendation evaluation. These limits are intentional: the goal of the present draft is to establish a clean tokenizer-side story before enlarging the experimental scope.

The next step is therefore straightforward. We will use the current best hierarchy-aware tokenizer as the new SID backbone for downstream SFT experiments, and test whether lower local ambiguity at the tokenizer stage translates into fewer same-prefix recommendation errors and better end-to-end recommendation quality.

Reproducibility note. All experiments in this draft are built as incremental modules on top of the existing OneRec repository. The default baseline path remains unchanged, and the graph-aware tokenizer branch uses isolated configs, scripts, logs, and outputs.

REFERENCES

- Dengzhao Fang, Jingtong Gao, Yu Li, Xiangyu Zhao, and Yi Chang. Prism: Purified representation and integrated semantic modeling for generative sequential recommendation, 2026. URL <https://arxiv.org/abs/2601.16556>.
- Xiaoyu Kong, Leheng Sheng, Junfei Tan, Yuxin Chen, Jiancan Wu, An Zhang, Xiang Wang, and Xiangnan He. Minionerec: An open-source framework for scaling generative recommendation, 2025. URL <https://arxiv.org/abs/2510.24431>.
- Doyup Lee, Chiheon Kim, Saehoon Kim, Minsu Cho, and Wook-Shin Han. Autoregressive image generation using residual quantization. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 11513–11522. IEEE, 2022. doi: 10.1109/CVPR52688.2022.01123. URL <https://doi.org/10.1109/CVPR52688.2022.01123>.
- Yu Liang, Zhongjin Zhang, Yuxuan Zhu, Kerui Zhang, Zhiluohan Guo, Wenhong Zhou, Zonqi Yang, Kangle Wu, Yabo Ni, Anxiang Zeng, Cong Fu, Jianxin Wang, and Jiazhi Xia. Rethinking generative recommender tokenizer: Recsys-native encoding and semantic quantization beyond llms, 2026. URL <https://arxiv.org/abs/2602.02338>.
- Chang Liu, Yimeng Bai, Xiaoyan Zhao, Yang Zhang, Fuli Feng, and Wenge Rong. Discrec: Disentangled semantic-collaborative modeling for generative recommendation, 2025. URL <https://arxiv.org/abs/2506.15576>.
- Enze Liu, Bowen Zheng, Cheng Ling, Lantao Hu, Han Li, and Wayne Xin Zhao. Generative recommender with end-to-end learnable item tokenization, 2024. URL <https://arxiv.org/abs/2409.05546>.
- Ahmad Bin Rabiah and Julian McAuley. Gsprec: Temporal-aware graph spectral filtering for recommendation, 2025. URL <https://arxiv.org/abs/2505.11552>.
- Shashank Rajput, Nikhil Mehta, Anima Singh, Raghunandan Hulikal Keshavan, Trung Vu, Lukasz Heldt, Lichan Hong, Yi Tay, Vinh Q. Tran, Jonah Samost, Maciej Kula, Ed H. Chi, and Mahesh Sathiamoorthy. Recommender systems with generative retrieval. In *Advances in Neural Information Processing Systems 36 (NeurIPS)*, 2023. URL http://papers.nips.cc/paper_files/paper/2023/hash/20dcab0f14046a5c6b02b61da9f13229-Abstract-Conference.html.
- Huanjie Wang, Xinchun Luo, Honghui Bao, Zixing Zhang, Lejian Ren, Yunfan Wu, Hongwei Zhang, Liwei Guan, and Guang Chen. Pit: A dynamic personalized item tokenizer for end-to-end generative recommendation, 2026. URL <https://arxiv.org/abs/2602.08530>.
- Jiafeng Xia, Dongsheng Li, Hansu Gu, Tun Lu, Peng Zhang, Li Shang, and Ning Gu. Frequency-aware graph signal processing for collaborative filtering, 2024. URL <https://arxiv.org/abs/2402.08426>.
- Bowen Zheng, Yupeng Hou, Hongyu Lu, Yu Chen, Wayne Xin Zhao, Ming Chen, and Ji-Rong Wen. Adapting large language models by integrating collaborative semantics for recommendation, 2023. URL <https://arxiv.org/abs/2311.09049>.
- Jieming Zhu, Mengqun Jin, Qijiong Liu, Zexuan Qiu, Zhenhua Dong, and Xiu Li. Cost: Contrastive quantization based semantic tokenization for generative recommendation, 2024. URL <https://arxiv.org/abs/2404.14774>.
- Tianyu Zhu, Yansong Shi, Yuan Zhang, Yihong Wu, Fengran Mo, and Jian-Yun Nie. Collaboration and transition: Distilling item transitions into multi-query self-attention for sequential recommendation, 2023. URL <https://arxiv.org/abs/2311.01056>.

A APPENDIX

A.1 CURRENT IMPLEMENTATION DETAILS

The current MGR-SID v1 implementation lives in an isolated experimental branch of the repository and does not alter the default MiniOneRec tokenizer path. The main training and generation entry points are:

- `scripts/experiment_mgr_sid_v1_train.py`
- `scripts/experiment_mgr_sid_v1_generate.py`

The core implementation is in:

- `src/onerec/experiments/mgr_sid/train_v1.py`
- `src/onerec/experiments/mgr_sid/transplanted_graph_bank.py`
- `src/onerec/experiments/mgr_sid/paper_transplants.py`

A.2 GRAPH CONSTRUCTION CHOICES IN V1

Table 5: Concrete graph-bank settings used by the current v1 experiments.

Component	Setting
History window	10
Coarse minimum edge weight	2.0
Local minimum edge weight	1.0
Semantic anchor top- k	32
Semantic anchor mixing weight	0.35
Spectral rank	48
Middle-band eigen ratio	[0.25, 0.65]
Temporal mixing ratio (alt. middle view)	0.35
Graph top- k per row in training	32

A.3 TOKENIZER TRAINING HYPERPARAMETERS

Table 6: Upstream-aligned tokenizer training regime used for the fair comparison in Table 3.

Hyperparameter	Value
Number of quantizers	3
Codebook sizes	[256, 256, 256]
Quantized embedding dimension	32
Encoder/decoder widths	[2048, 1024, 512, 256, 128, 64]
Optimizer	AdamW
Learning rate	10^{-3}
Weight decay	0
Scheduler	constant with warmup
Warmup epochs	50
Epochs	10000
Batch size	20480
Quantization loss weight	1.0
β	0.25
Graph loss weights	(0.05, 0.15, 0.05)

A.4 ADDITIONAL INTERPRETATION OF THE LOCAL AMBIGUITY RESULT

The local ambiguity analysis compares the final baseline index with the final hierarchy-aware index rather than comparing recommendation outputs. This makes the analysis tokenizer-native. The result is therefore best interpreted as evidence about the structure of the tokenizer itself:

- the hierarchy-aware tokenizer produces less crowded same- l_2 buckets;
- the reduction is visible both at catalog level and when weighted by test targets;
- the structural change is aligned with the motivating failure mode of local leaf ambiguity.

A.5 SCOPE OF THE CURRENT DRAFT

This paper should be read as a tokenizer-stage draft rather than a finished end-to-end paper. The remaining milestones are:

1. add repeated-seed verification for the final tokenizer comparison;
2. evaluate the hierarchy-aware tokenizer on Office under the same upstream-aligned regime;
3. feed the improved SID into SFT and then RL;
4. compare against more published tokenizer baselines once downstream metrics are available.